



COMPLIANCE OPERATIONS BRIEFING

SEPTEMBER 2026

The transaction intelligence your TMS is missing.

Transaction monitoring intelligence, robotic process automation and agentic AI — deployed on top of the systems your bank already runs. No replacement. No migration. No downtime.

TMS-AGNOSTIC

CBS-AGNOSTIC

API-FIRST

–40%

reduction in total alert volume

LIVE BANK ENGAGEMENT

–28%

reduction in false positives

VALIDATED ON THEIR OWN HISTORY

9

automations built, governed and deployable

RPA AND AGENTIC AI

Five questions, answered in order.

What is built, what is measured, and what it takes to start.

01

The problem

Where compliance effort is actually spent — and why 70%+ of Level-1 capacity is consumed before investigation begins.

02

The window

Five regulators moving in one direction, and what each of them now expects of a monitoring programme.

03

The platform

Compliance Catalyst — four layers that unify the data, simulate the rules, automate the queue and govern the lot.

04

What we built

Nine automations across RPA and agentic AI, each with a named control position.

05

The proof

Measured results from a live bank engagement, and a 90-day path from pilot to production.

Optimoney at a glance

AML and Fraud operations tech for banks — built by bankers, chartered accountants and engineers who have worked inside financial services.

2023

Founded; building for Indian banking since

4

Layers — data, simulation, automation, governance

9

Automations built, governed and deployable today

9

Advisors from RBI, HDFC Bank, SBI and industry

What we do

- AML and fraud operations technology for Indian banks
- Compliance Catalyst — four-layer AML overlay
- RegTech, analytics and payments — API-driven

How we integrate

- Sits above the existing TMS, fraud engine and core
- No replacement, no migration, no downtime
- Hybrid cloud or on-premise deployment

Where we work

- Live with a systemically important Indian private bank
- Cross-border AML mandate with a Gulf banking operation
- Scheduled, small finance and co-operative banks

How we are governed

- Advisory council of former regulators and bank officers
- Knowledge partner to NIBM, Pune
- Built to FIU-IND, RBI and FATF frameworks

Client identities withheld by agreement.

AML teams are drowning in their own alerts.

95

out of every 100 AML alerts
are false positives.

Analyst burnout

80%+ of the analyst's day goes to clerical work, not judgment. Real investigation gets what is left.

Real threats slip through


High-risk transactions are buried in the alert flood. The signal is lost before it can be acted on.

Turnaround breaches

14–21 days from alert to STR against a tightening regulatory clock.

Cost of noise


₹800–2,000 per false positive — a tax that compounds every single month.



5 in 100 · the alerts that actually deserved an analyst

Where compliance effort is actually spent.

RELATIVE SHARE OF COMPLIANCE-TEAM EFFORT, BY LIFECYCLE STAGE



● Clerical work

Gathering, re-keying, chasing, assembling, logging.

→ Addressed by RPA

● Judgment work

Tiering, segmentation, investigation, narrative.

→ Assisted by agentic AI

● Detection

The rule engine firing an alert.

→ Already automated

The cost of doing nothing is no longer theoretical.

Money laundering is not a compliance problem. It is a structural failure — and supervisors have stopped treating it as a filing exercise.

GLOBAL

\$2T+

laundered globally each year

2–5% of world GDP

~1%

of illicit flows ever detected

99¢ in every dollar escapes

\$4.5B

in global AML fines in 2024

72% traced to monitoring failures

IN THE PRESS

A Decade-Long Investor Nightmare

PACL Money Laundering Case ED Attaches Properties Rs 696 Crore

Web of Front Companies Used to Launder Rs 48,000 Crore

Enforcement Directorate's latest crackdown in the Rs 48,000 crore PACL Ponzi scam. Rs 696 crore Panchkula assets attached, revealing complex layering by Bhangoo-led group. A Sprouts News Investigation.

INDIA

₹36,014 Cr

bank fraud in FY25

3× the prior year

4.5 lakh

mule accounts frozen

₹17,000 Cr siphoned first

88%

rise in RBI enforcement

2021–2024

₹54.78 Cr

RBI penalties in FY25


353 entities penalised

TD Bank paid \$3B for precisely the failure this platform addresses — inadequate transaction monitoring and missed STR filings. India's supervisors are on the same trajectory.



BCs to PAs, and then to a Black Hole, Money Laundering Takes New Route

Payout APIs of banks being misused to move cash; industry alerts RBI, Financial Intelligence Unit



Sources: RBI Annual Report FY25 & FY26 · UNODC · Fenargo AML Enforcement 2024 · I4C / CFCFRS

Why tuning the same engine harder does not work.

01

Fragmented data

CBS, UPI, NEFT, RTGS and IMPS sit in separate silos. There is no unified customer risk view, and entity relationships stay invisible across channels.

Consequence: the same customer is assessed five times, five different ways.

02

Static, uncalibrated rules

Rules are recalibrated infrequently and never simulated. The same rule fires for a ₹50,000 transfer and a ₹5 crore one.

Consequence: thresholds are defended by habit, not by evidence.

03

Zero process governance

L1 → L2 → L3 handoffs run over email and Excel. There is no MLRO command centre, and the audit trail is stitched together after the fact.

Consequence: the inspection file is reconstructed, not retrieved.

Five mandates. One direction. The window is now.

Compliance has stopped being a back-office function. Every regulator with a view on financial crime is pulling the same way.

RBI

Customer at the pivot

Risk-based monitoring centred on the customer, not the transaction.

FIU-IND

Define your thresholds

Prescribed rules are a floor — banks must set and defend their own parameters.

I4C / MHA

Fuse fraud with AML

Financial crime and laundering brought under one PMLA lens.

NPCI

UPI / IMPS risk

Real-time monitoring across India's highest-volume payment rails.

DoT / MNRL

Risk-tag the number

Mobile-number risk signals feeding the unified customer view.

Supervisory posture

Inspections increasingly ask for evidence of tuning and testing, not just a rulebook.

Board accountability

AI and automation now sit inside a board-owned policy under RBI's FREE-AI framework.

Cost of delay

Every quarter without simulation is another quarter of thresholds you cannot defend.

Compliance Catalyst — four layers above your TMS.

Unify the data, simulate the rules, automate the queue and govern the lot — sitting on top of what you already run. No replacement, no downtime, no three-year programme.

L1 RECON

A single entity-linked data layer

CBS, UPI, NEFT, RTGS and IMPS reconciled into one clean view of the customer. Entity relationships become visible across every channel.

Eliminates the root cause of most false positives.

L2 SIMULATION

Backtest before you deploy

Every rule modification is simulated against real historical data first. An agentic assistant helps configure and tune rules at the root — not downstream after the noise lands.

Know the false-positive rate before a single analyst sees it.

L3 AUTOMATION

Alerts that route themselves

Agentic AI smart-triages every alert — scoring contextual risk, drafting case summaries and routing to the right analyst tier. L1 → L2 → L3 handoffs are automated and tracked.

SLA-bound from the moment an alert is born.

L4 GOVERNANCE

The compliance officer's cockpit

Real-time visibility into every pending case, officer workload, STR-pipeline integrity and turnaround metric. The MLRO commands, instead of reconstructing.

AI-drafted case summaries keep it audit-ready, always.

Built to integrate — not to replace.

Nothing is torn out. We read from the rails you already run, and write into the frameworks you already report against.

No rip-and-replace

Integrate alongside what you run today — nothing gets torn out.


Core-agnostic

No CBS lock-in to Finacle, Temenos or any specific core.

Zero-downtime

Live transition with a 99.99% uptime SLA — no disruption to operations.

YOUR SOURCE SYSTEMS



Nine automations, in three groups.

RPA · 01-05

Removing clerical effort

Deterministic bots, scripted to your procedure.

- 01 Alert data aggregation
- 02 KYC refresh & periodic review
- 03 Regulatory report assembly
- 04 Cross-rail reconciliation
- 05 Level-1 screening dispositions

AGENTIC AI · 06-09

Supporting judgment

Models and agents that rank, cluster and draft.

- 06 Smart risk tiering & triage
- 07 Dynamic customer segmentation
- 08 Investigation & STR co-pilot
- 09 Entity resolution & networks

CONTROLS · ACROSS ALL

Keeping both defensible

Built in from day one, not added at inspection.

- Maker-checker at every tier
- No automated closure of alerts
- Explainable scores, reason codes
- Append-only audit trail
- Named owner for every bot

Robotic process automation — five use cases.

Deterministic bots scripted to your written procedure. They gather, chase, assemble and log. They do not decide.

#	AUTOMATION	WHAT THE BOT DOES	MANUAL EFFORT REMOVED
01	Alert data aggregation	Pulls KYC, 12-month behaviour, prior alerts and screening hits into one case pack the moment an alert is raised.	Analyst search across four to six systems, per alert.
02	KYC refresh & periodic review	Scans for due and overdue reviews, chases documents with escalation, updates records and routes exceptions.	Branch follow-up, tracking registers, ageing reports.
03	Regulatory report assembly	Extracts prescribed fields, validates against format rules, assembles annexures and tracks acknowledgements.	Manual compilation and format rework before filing.
04	Cross-rail reconciliation	Matches CBS, UPI, NEFT, RTGS and IMPS entries continuously and flags breaks to a named owner.	Overnight batch reconciliation and break chasing.
05	Level-1 screening dispositions	Clears unambiguous mismatches against documented criteria and links repeat hits to prior dispositions.	Repeat clearing of the same obvious name-screening hits.

Agentic AI — four deployed use cases.

Models that rank, cluster and draft — each one simulated on your own history before a single threshold moves.

#	AUTOMATION	WHAT THE MODEL DOES	WHERE THE HUMAN STAYS
06	Smart risk tiering & triage	Scores every alert on likelihood of being genuinely suspicious, then enriches and routes by tier.	Every disposition, including low-scored alerts under sampling.
07	Dynamic customer segmentation	Clusters customers on observed behaviour and sets thresholds per micro-segment.	Approval of any threshold change, after below-the-line testing.
08	Investigation & STR co-pilot	Assembles evidence, reconstructs the money trail, drafts a cited narrative in filing format.	Review, edit, approval and filing — maker–checker intact.
09	Entity resolution & networks	Links identities across rails and surfaces mule clusters and layering chains as one case.	Any freeze, exit or customer action.

Built to withstand an inspection.

Automation carries its own risk register. Bots and models are governed from day one, in line with RBI's FREE-AI framework and existing IT and outsourcing norms.

PROCESS AUTOMATION CONTROLS

Ownership and inventory

Every bot has a named owner, a documented scope and an exit process. No undocumented automation.

Identity and access

Bot credentials sit in the password vault and within periodic access reviews and IS audit scope.

Failure and continuity

Breaks are alerted same-day, change logs are versioned, and the bot estate sits inside the BCP.

MODEL CONTROLS

Board-owned AI policy

Approved policy, defined accountability and human oversight for consequential decisions.

Explainability and validation

Scores decomposed into driving features. Independent validation, drift monitoring, periodic recalibration.

Human primacy

AI recommends, ranks and drafts. Officers close, file and act. No exception in the reporting pipeline.

From alert fatigue to defensible throughput.

What changes for a compliance operation once the layer above the engine is in place.

BEFORE	METRIC	WITH OPTIMONEY
95%	False positive rate	55%
14–21 days	Alert-to-STR cycle time	10–12 days
2–5 months per rule change	Rule tuning	AI-assisted self-serve weeks, simulated first
Manual write-ups hours per case	Case summaries & triage	Agentic AI drafts minutes, analyst-reviewed
Excel + informal	MLRO visibility	Real-time command centre
Manual reconstruction	Regulatory audit readiness	One-click audit trail

Ranges reflect the live engagement and comparable industry benchmarks. Every figure is re-established on your own data during the pilot.

Measured on a live bank engagement.

A systemically important Indian private bank with one of the largest AML compliance teams in the country. Identity withheld by agreement.

Starting position

- 7,000 alerts and 100+ RFIs per month
- Rule engine generating high false-positive volume
- Heavy manual reconciliation between systems

What we implemented

- Aggregated data into a single customer view
- Mirrored existing rules in a simulation engine
- Optimised thresholds against their own history

How it was validated

- Replayed on historical data before deployment
- Below-the-line testing on every threshold change
- Detection coverage confirmed unchanged

-40%

reduction in total alert volume

-28%

reduction in false positives

+50%

improvement in STR filing quality

Zero

detection coverage compromised

What made it defensible: every threshold change was replayed on the bank's own history and evidenced with below-the-line testing before it went anywhere near production.

A 90-day deployment, in four stages.

Process automation and model work run in parallel. Each stage closes with a defined deliverable.

DAYS 0–15

Discovery and prioritisation

Walk the alert lifecycle, quantify effort per stage, select the three highest-volume, lowest-ambiguity processes. **Deliverable: costed automation map.**

DAYS 16–45

First three bots in production

Built, tested, credentialed, documented and handed over with named owners. **Deliverable: measured effort saved, plus bot audit pack.**

DAYS 46–75

Model layer simulated on your history

Risk tiering and segmentation replayed on 6–12 months of your data, with no production exposure. **Deliverable: projected FP reduction with BTL results.**

DAYS 76–90

Co-pilot live within maker–checker

Case assembly and drafting for one alert type; analyst approves, Principal Officer signs. **Deliverable: cycle time, before and after.**

WHAT MAKES THIS PACE POSSIBLE

No system replacement. API-first integration alongside the existing TMS and core.

Simulation before deployment. Impact is evidenced on history, not gated on production trials.

Banking-native team. No ramp-up on AML process or regulatory expectation.

Parallel workstreams. Bots, models and governance are built concurrently.

STARTING POINT

A two-week proof of concept on a representative sample from your systems, at no commitment.

Built by people who have run these systems.

Chartered accountants, bankers and engineers from inside financial services — not a generalist software team that arrived at compliance — backed by former supervisors and principal officers.



Mihir Mehta

CO-FOUNDER & CEO



Sushil Kurri

CO-FOUNDER & CTO



Rajat A

DIRECTOR



Rohit Vaidya

HEAD OF PRODUCTS



Bibek Kartik

OPERATIONS

ADVISORY COUNCIL

Uma Shankar Paliwal · Former Executive Director, RBI

Neelkanth Chatterjee · Ex ICICI Lombard, Citi, SCB

Vivek Gupta · Transaction Square

Ravi Lahoty · Ex Principal Officer, HDFC Bank India

Nirav Mehta · MD, Morgan Stanley PE India

Karan Khetrapal · CEO, Indian Express CV

Hemant Adarkar · Independent Director, SBI

Suneel Aiyar · Former CEO, Writer Corporation

Venkat Hariharan · Open Invention Network



The noise was the problem. We fixed the noise.

Give us a representative transaction sample and two weeks. We will show you the false-positive reduction on your own data — before you commit to anything.

Mihir Mehta

Co-founder & CEO, Optimoney Technologies

mihirmehta@optimoneytech.com · optimoneytech.com

THE TWO-WEEK PROOF OF CONCEPT

What we need

A representative transaction sample and your current rule set. Nothing touches production.

What we do

Mirror your rules in the simulation engine and replay them against your own history.

What you get

A measured false-positive reduction, a costed automation map, and below-the-line test evidence.

What it costs

Nothing. No commitment until the simulation speaks for itself.